

Algorithms for Image Analysis

Elements of Image (Pre)-Processing and Feature Detection

Acknowledgements: slides from Steven Seitz, Aleosha Efros, David Forsyth, and Gonzalez & Woods

Image Processing Basics

■ Point Processing

- gamma correction
- window-center correction
- histogram equalization

Extra Reading: Szeliski, Sec 3.1

intensities, colors

■ Filtering (linear and non-linear)

- mean, Gaussian, and median filters
- image gradients, Laplacian
- normalized cross-correlation (NCC)
- etc...: Fourier, Gabor, wavelets (Szeliski, Sec 3.4-3.5)

Extra Reading: Szeliski, Sec 3.2-3.3

contrast edges

texture

templates, patches

■ Other features

Extra Reading: Szeliski, Sec. 4.1

Harris corners, MOPS, SIFT, etc.

Summary of image transformations

- An **image processing** operation (or transformation) typically defines a new image g in terms of an existing image f .

Examples:

Summary of image transformations

- An **image processing** operation (or transformation) typically defines a new image g in terms of an existing image f .

Examples:

- **Geometric (domain) transformation:**

- What kinds of operations $g(x, y) = f(t_x(x, y), t_y(x, y))$

Summary of image transformations

- An **image processing** operation (or transformation) typically defines a new image g in terms of an existing image f .

Examples:

- **Geometric (domain) transformation:**

- What kinds of operations $g(x, y) = f(t_x(x, y), t_y(x, y))$

- **Range transformation:**

- What kinds of operations $g(x, y) = t(f(x, y))$

Summary of image transformations

- An **image processing** operation (or transformation) typically defines a new image g in terms of an existing image f .

Examples:

- **Geometric (domain) transformation:**

- What kinds of operations $g(x, y) = f(t_x(x, y), t_y(x, y))$

- **Range transformation:**

- What kinds of operations $g(x, y) = t(f(x, y))$

point processing

- **Filtering** also generates new images from an existing image

neighborhood
processing

- more on filtering later

$$g(x, y) = \int_{\substack{|u| < \varepsilon \\ |v| < \varepsilon}} h(u, v) \cdot f(x - u, y - v) \cdot du \cdot dv$$

Point Processing

$$g(x, y) = t(f(x, y))$$

for each original image intensity value I function $t(\cdot)$
returns a transformed intensity value $t(I)$.

$$I' = t(I)$$

NOTE: we will often use
notation I_p instead of $f(x, y)$ to
denote intensity at pixel $p=(x, y)$

- **Important:** every pixel is for itself
 - spatial information is ignored!
- What can point processing do?


(we will focus on grey scale images, see Szeliski 3.1 for examples of point processing for color images)

Point Processing:

Examples of gray-scale transforms t


$$I' = t(I)$$

FIGURE 3.3 Some basic gray-level transformation functions used for image enhancement.




Point Processing:

Negative



I_p or $f(x, y)$



I'_p or $g(x, y)$

a b

FIGURE 3.4
(a) Original digital mammogram.
(b) Negative image obtained using the negative transformation in Eq. (3.2-1).
(Courtesy of G.E. Medical Systems.)

$$t(I) = 255 - I$$

$$g(x, y) = t(f(x, y)) = 255 - f(x, y)$$

Point Processing:

Power-law transformations t




FIGURE 3.6 Plots of the equation $s = cr^\gamma$ for various values of γ ($c = 1$ in all cases).

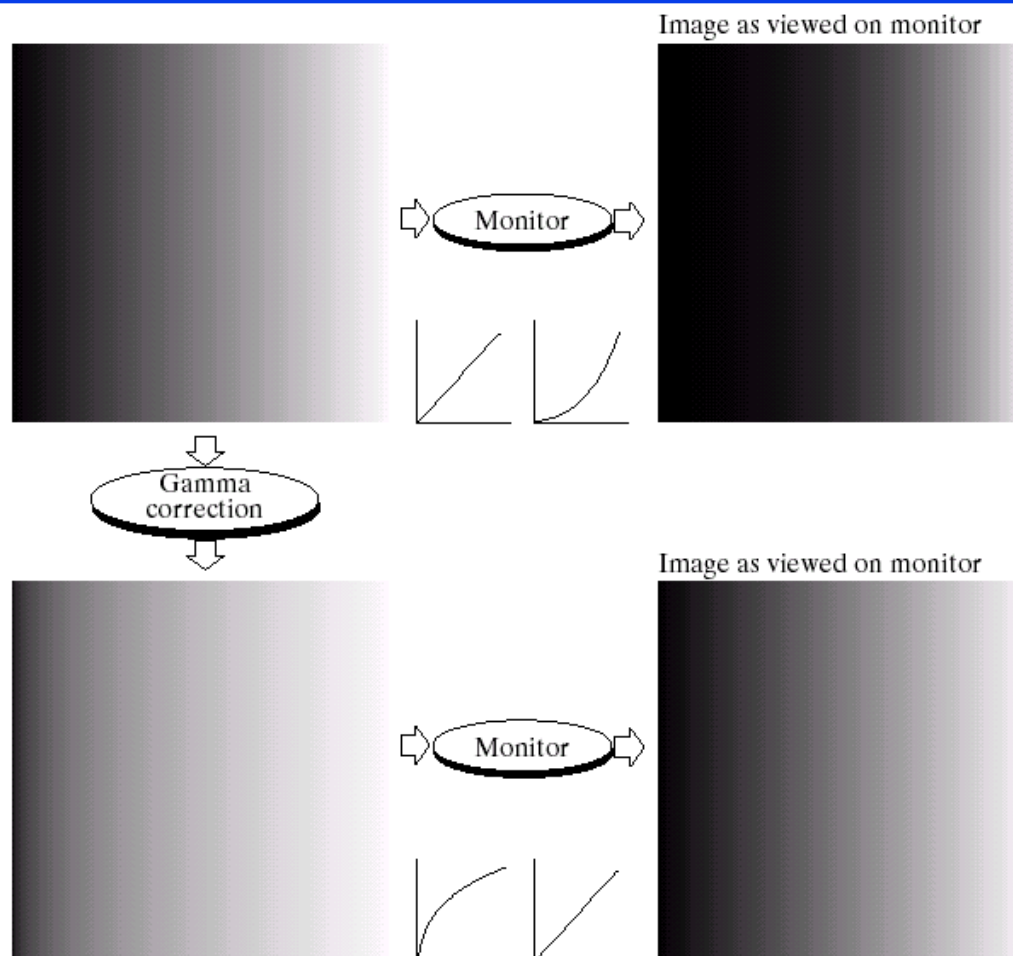
Point Processing:

Gamma Correction

a	b
c	d

FIGURE 3.7

(a) Linear-wedge gray-scale image.
(b) Response of monitor to linear wedge.
(c) Gamma-corrected wedge.
(d) Output of monitor.



Gamma Measuring Applet:

<http://www.cs.berkeley.edu/~efros/java/gamma/gamma.html>



Point Processing:

Enhancing Image via Gamma Correction

a	b
c	d

FIGURE 3.9

(a) Aerial image.
(b)–(d) Results of
applying the
transformation in
Eq. (3.2-3) with
 $c = 1$ and
 $\gamma = 3.0, 4.0,$ and
 5.0 , respectively.
(Original image
for this example
courtesy of
NASA.)



Point Processing:

Understanding Image Histograms

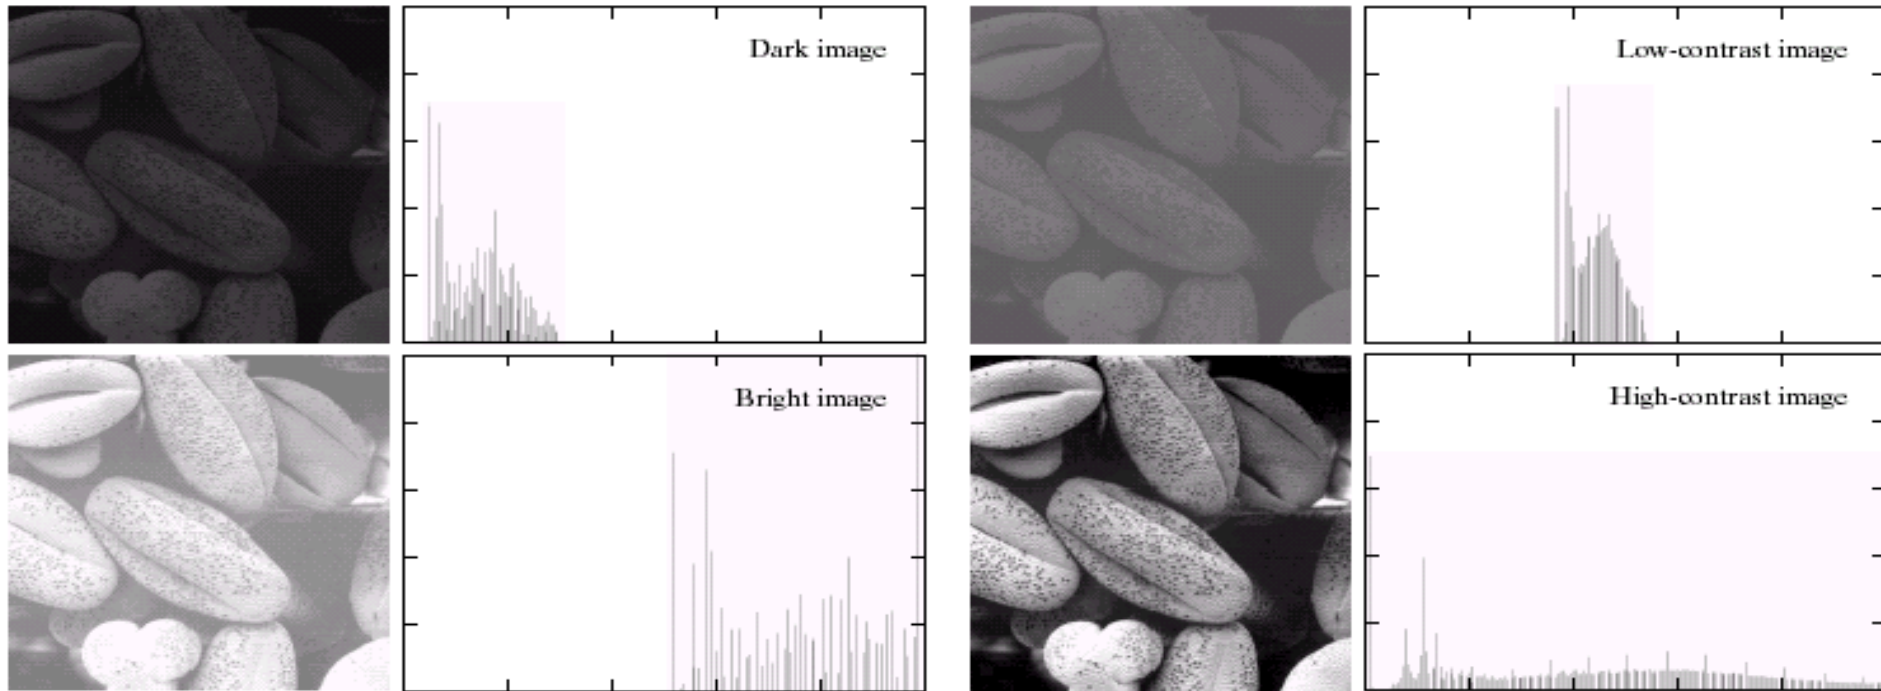





Image Brightness

Image Contrast

probability of intensity i : $p(i) = \frac{n_i}{n}$ ---number of pixels with intensity i
 ---total number of pixels in the image

Point Processing:

Contrast Stretching



a	b
c	d

FIGURE 3.10

Contrast stretching.

(a) Form of transformation function. (b) A low-contrast image. (c) Result of contrast stretching. (d) Result of thresholding.

(Original image courtesy of Dr. Roger Heady, Research School of Biological Sciences, Australian National University, Canberra, Australia.)


Point Processing:

Contrast Stretching

Original images

Histogram corrected images

1)



2)





Point Processing:

Contrast Stretching


Original images

Histogram corrected images

3)



4)




One way to automatically select transformation t :

Histogram Equalization

FIGURE 3.18

Transformation functions (1) through (4) were obtained from the histograms of the images in Fig.3.17(a), using Eq. (3.3-8).



$$t(i) = \sum_{j=0}^i p(j) = \sum_{j=0}^i \frac{n_j}{n} \quad = \text{cumulative distribution of image intensities}$$

...see Gonzalez and Woods, Sec3.3.1, for more details

Point processing

Histogram Equalization

$$t(i) = \sum_{j=0}^i p(j) = \sum_{j=0}^i \frac{n_j}{n} \quad = \text{cumulative distribution of image intensities}$$

Why does that work?

Answer in probability theory:

I – random variable with *probability* distribution $p(i)$ over i in $[0,1]$


If $t(i)$ is a *cumulative* distribution of I then

$I' = t(I)$ – is a random variable with *uniform* distribution over its range $[0,1]$

That is, transform image I' will have a uniformly-spread histogram (good contrast)


Point Processing:

Window-Center adjustment




Point Processing:

Window-Center adjustment



Point Processing:

Window-Center adjustment




Point Processing:

Window-Center adjustment

Window = 4000

Center = 500




Point Processing:

Window-Center adjustment

Window = 2000

Center = 500




Point Processing:

Window-Center adjustment

Window = 800


Center = 500



Point Processing:

Window-Center adjustment

Window = 0
Center = 500




If $window=0$ then we get
binary image *thresholding*

Point Processing:

Window-Center adjustment




Window = 800
Center = 500



Window = 800
Center = 1160

Neighborhood Processing (or filtering)

- Q: What happens if I reshuffle all pixels within the image?




- A: It's histogram won't change.
No point processing will be affected...
- Images contain a lot of “spatial information”

Readings: [Szeliski, Sec 3.2-3.3](#)

Neighborhood Processing (filtering)

Linear image transforms

Let's start with 1D image (a signal): $f[i]$



A very general and useful class of transforms are the **linear transforms** of f , defined by a matrix M

$$\begin{bmatrix} * & * & \dots & * \\ * & * & \dots & * \\ \vdots & \vdots & \ddots & \vdots \\ * & * & \dots & * \end{bmatrix}
 \begin{bmatrix} * \\ * \\ \vdots \\ * \end{bmatrix}
 =
 \begin{bmatrix} * \\ * \\ \vdots \\ * \end{bmatrix}$$


$M[i, j]$ $f[i]$ $g[i]$

$$g[i] = \sum_{j=1} M[i, j] f[j]$$

Neighborhood Processing (filtering)

Linear image transforms

Let's start with 1D image (a signal): $f[i]$



$$f[i] = \begin{bmatrix} \\ \\ \\ \\ \\ \\ \\ \\ \end{bmatrix}$$

matrix M

$$\begin{bmatrix} 1 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 1 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 1 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 1 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 1 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 1 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 1 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 & 1 \end{bmatrix}$$

$$f[i] \rightarrow$$


$$\begin{bmatrix} 1 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 1 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 1 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 1 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 1 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 1 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 1 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 1 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \end{bmatrix}$$

$$f[i] \rightarrow$$

Neighborhood Processing (filtering)

Linear image transforms

Let's start with 1D image (a signal): $f[i]$



$$f[i] = \begin{bmatrix} 4 \\ 4 \\ 1 \\ 1 \\ 3 \\ 5 \\ 5 \\ 5 \\ 5 \end{bmatrix}$$

matrix M

$$\begin{bmatrix} 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 1 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 1 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 1 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 1 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 1 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 1 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 1 & 0 \end{bmatrix}$$

$$f[i] \rightarrow$$

$$\frac{1}{2} \begin{bmatrix} 2 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 1 & 1 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 1 & 1 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 1 & 1 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 1 & 1 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 1 & 1 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 1 & 1 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 1 & 1 \end{bmatrix}$$

$$f[i] \rightarrow$$

Neighborhood Processing (filtering)

Linear shift-invariant filters

matrix M

$$\begin{bmatrix} * & * & 0 & 0 & 0 & 0 & 0 & 0 \\ a & b & c & 0 & 0 & 0 & 0 & 0 \\ 0 & a & b & c & 0 & 0 & 0 & 0 \\ 0 & 0 & a & b & c & 0 & 0 & 0 \\ 0 & 0 & 0 & a & b & c & 0 & 0 \\ 0 & 0 & 0 & 0 & a & b & c & 0 \\ 0 & 0 & 0 & 0 & 0 & a & b & c \\ 0 & 0 & 0 & 0 & 0 & 0 & * & * \end{bmatrix}$$

This pattern is very common

- same entries in each row
- all non-zero entries near the diagonal

$$g = M \cdot f$$

It is known as a **linear shift-invariant filter** and is represented by a **kernel** (or **mask**) h :

$$h[i] = [a \ b \ c]$$

and can be written (for kernel of size $2k+1$) as:

$$g[i] = \sum_{u=-k}^k h[u] \cdot f[i+u]$$

The above allows negative filter indices. When you implement need to use: $h[u+k]$ instead of $h[u]$

Neighborhood Processing (filtering)

2D linear transforms

We can do the same thing for 2D images by concatenating all of the rows into one long vector (in a “*raster-scan*” order):

$$\hat{f}[i] = f[\lfloor i/m \rfloor, i \% m]$$

$$\begin{array}{ccc}
 \begin{bmatrix} * & * & \dots & * \\ * & * & \dots & * \\ \vdots & \vdots & \ddots & \vdots \\ * & * & \dots & * \end{bmatrix} & \begin{bmatrix} * \\ * \\ \vdots \\ * \end{bmatrix} & = & \begin{bmatrix} * \\ * \\ \vdots \\ * \end{bmatrix} \\
 M[i, j] & \hat{f}[i] & & \hat{g}[i]
 \end{array}$$

Neighborhood Processing (filtering)

2D filtering

A 2D image $f[i,j]$ can be filtered by a 2D kernel $h[u,v]$ to produce an output image $g[i,j]$:

$$g[i,j] = \sum_{u=-k}^k \sum_{v=-k}^k h[u,v] \cdot f[i+u, j+v]$$

This is called a **cross-correlation** operation and written:

$$g = h \circ f$$

h is called the “**filter**,” “**kernel**,” or “**mask**.”

Neighborhood Processing (filtering)

2D filtering

A **convolution** operation is a cross-correlation where the filter is flipped both horizontally and vertically before being applied to the image:

$$g[i, j] = \sum_{u=-k}^k \sum_{v=-k}^k h[u, v] \cdot f[i - u, j - v]$$

It is written: $g = h * f$ $= \sum_{u=-k}^k \sum_{v=-k}^k h[-u, -v] \cdot f[i + u, j + v]$

How does convolution differ from cross-correlation?

If $h[u, v] = h[-u, -v]$ then there is no difference between convolution and cross-correlation

convolution has additional “technical” properties: **commutativity**, **associativity**. Also, “nice” properties wrt Fourier analysis.
(see Szeliski Sec 3.2, Gonzalez and Woods Sec. 4.6.4)

2D filtering

Noise

Filtering is useful for
noise reduction...

(side effects: **blurring**)

Common types of noise:

- **Salt and pepper noise:** random occurrences of black and white pixels
- **Impulse noise:** random occurrences of white pixels
- **Gaussian noise:** variations in intensity drawn from a Gaussian normal distribution



Original



Salt and pepper noise



Impulse noise



Gaussian noise

[illegible]

Neighborhood Processing (filtering)

Mean filtering

0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	0	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0

 $F[x, y]$

 $G[x, y]$

Neighborhood Processing (filtering)

Mean filtering

0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	0	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0

 $F[x, y]$

	0	10	20	30	30	30	20	10	
	0	20	40	60	60	60	40	20	
	0	30	60	90	90	90	60	30	
	0	30	50	80	80	90	60	30	
	0	30	50	80	80	90	60	30	
	0	20	30	50	50	60	40	20	
	10	20	30	30	30	30	20	10	
	10	10	10	0	0	0	0	0	

 $G[x, y]$

Effect of mean filters

Gaussian
noise

Salt and pepper
noise

3x3



5x5



7x7




Neighborhood Processing (filtering)

Mean kernel

- What's the kernel for a 3x3 mean filter?

0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	0	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0

 $F[x, y]$  $H[u, v]$

Neighborhood Processing (filtering)


Gaussian Filtering

- A Gaussian kernel gives less weight to pixels further from the center of the window

0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	0	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0

 $F[x, y]$

$$\frac{1}{16} \cdot \begin{bmatrix} 1 & 2 & 1 \\ 2 & 4 & 2 \\ 1 & 2 & 1 \end{bmatrix} H[u, v]$$





$$h(u, v) = \frac{1}{2\pi\sigma^2} e^{-\frac{u^2+v^2}{\sigma^2}}$$

This kernel is an approximation of a Gaussian function:

Neighborhood Processing (filtering)

Mean vs. Gaussian filtering



Median filters

- A **Median Filter** operates over a window by selecting the median intensity in the window.
- What advantage does a median filter have over a mean filter?
- Is a median filter a kind of convolution?
 - No, median filter is an example of non-linear filtering

Comparison: salt and pepper noise

3x3

Mean



Gaussian



Median



5x5




7x7



Comparison: Gaussian noise

3x3

Mean



Gaussian




Median



5x5



7x7



Reading: Forsyth & Ponce, 8.1-8.2

Differentiation and convolution

- Recall

$$\frac{\partial}{\partial x} f = \lim_{\varepsilon \rightarrow 0} \left(\frac{f(x + \varepsilon, y) - f(x, y)}{\varepsilon} \right)$$

- Now this is linear and shift invariant, so must be the result of a convolution.

- We could approximate this as

$$\frac{\partial}{\partial x} f \approx \frac{f(x_{i+1}, y) - f(x_{i-1}, y)}{2 \cdot \Delta x}$$

$$= \nabla_x * f \quad (\text{convolution})$$

with kernel $\frac{1}{2\Delta x} \cdot$

0	0	0
1	0	-1
0	0	0

$$\nabla_x [u, v]$$

sometimes this may not be a very good way to do things, as we shall see

Reading: Forsyth & Ponce, 8.1-8.2

Differentiation and convolution

- Recall

$$\frac{\partial}{\partial x} f = \lim_{\varepsilon \rightarrow 0} \left(\frac{f(x + \varepsilon, y) - f(x, y)}{\varepsilon} \right)$$

- Now this is linear and shift invariant, so must be the result of a convolution.

- We could approximate this as

$$\frac{\partial}{\partial x} f \approx \frac{f(x_{i+1}, y) - f(x_{i-1}, y)}{2 \cdot \Delta x}$$

$$= \nabla_x * f \quad (\text{convolution})$$

with kernel $\frac{1}{2\Delta x} \cdot$


0	0	0
1	0	-1
0	0	0

$$\nabla_x [u, v]$$


sometimes this may not be a very good way to do things, as we shall see



Finite differences



Finite differences responding to noise





Increasing noise ->
(this is zero mean additive gaussian noise)

Finite differences and noise


- Finite difference filters respond strongly to noise
 - obvious reason: image noise results in pixels that look very different from their neighbours
- Generally, the larger the noise the stronger the response
- What is to be done?
 - intuitively, most pixels in images look quite a lot like their neighbours
 - this is true even at an edge; along the edge they're similar, across the edge they're not
 - suggests that smoothing the image should help, by forcing pixels different to their neighbours (=noise pixels?) to look more like neighbours

Smoothing and Differentiation


- Issue: noise
 - smooth before differentiation
 - two convolutions: smooth, and then differentiate?
 - actually, no - we can use a derivative of Gaussian filter
 - because differentiation is convolution, and convolution is associative $\nabla_x * (H * f) = (\nabla_x * H) * f$

 $\nabla_x * H$  $\nabla_y * H$ 


$$(\nabla_x * H) * f$$



1 pixel



3 pixels



7 pixels

The scale of the smoothing filter affects derivative estimates, and also the semantics of the edges recovered.

Sobel derivative kernels

$$\frac{\partial}{\partial x} f$$

$$\frac{\partial}{\partial y} f$$

$$\frac{1}{8\Delta x} \cdot \begin{array}{|c|c|c|} \hline 1 & 0 & -1 \\ \hline 2 & 0 & -2 \\ \hline 1 & 0 & -1 \\ \hline \end{array}$$

$$\nabla_x[u, v]$$

$$\frac{1}{8\Delta y} \cdot \begin{array}{|c|c|c|} \hline 1 & 2 & 1 \\ \hline 0 & 0 & 0 \\ \hline -1 & -2 & -1 \\ \hline \end{array}$$

$$\nabla_y[u, v]$$


Image Gradients

- Recall for a function of two (or more) variables $f(x, y)$

**Gradient
at point (x,y)**

$$\nabla f = \begin{bmatrix} \frac{\partial f}{\partial x} \\ \frac{\partial f}{\partial y} \end{bmatrix} \approx \begin{bmatrix} \nabla_x * f \\ \nabla_y * f \end{bmatrix}$$

a two (or more)
dimensional vector



- The absolute value $|\nabla f| = \sqrt{\left(\frac{\partial f}{\partial x}\right)^2 + \left(\frac{\partial f}{\partial y}\right)^2} \approx \sqrt{(\nabla_x * f)^2 + (\nabla_y * f)^2}$ is large at image boundaries

small image
gradients in low
textured areas

- The direction of the gradient corresponds to the direction of the “steepest ascend”
- normally gradient is orthogonal to object boundaries in the image.

Comment: vector ∇f is independent of specific coordinate system


- Equivalently, *gradient* of function $f(p)$ at point $p \in \mathbb{R}^2$ can be defined as a vector ∇f s.t. for any unit vector \vec{n}

**Gradient
at point p**

$$(\nabla f \cdot \vec{n}) = \frac{\partial f}{\partial \vec{n}} \approx \frac{f(p + \varepsilon \cdot \vec{n}) - f(p)}{\varepsilon}$$

dot product




directional derivative of
function f along direction \vec{n}



- pure vector algebra, specific coordinate system is irrelevant
- works for functions of two, three, or any larger number of variables
- previous slide gives a specific way for computing coordinates $(\frac{\partial f}{\partial x}, \frac{\partial f}{\partial y})$ of vector ∇f w.r.t. given orthogonal basis (axis X and Y).

Image Gradient

- Typical application where image gradients are used is *image edge detection*
 - find points with large image gradients



“edge features”

Canny edge detector suppresses
non-extrema Gradient points

Second Image Derivatives

(Laplace operator Δf)

$$\Delta f = \frac{\partial^2 f}{\partial x^2} + \frac{\partial^2 f}{\partial y^2} = \begin{bmatrix} \frac{\partial}{\partial x} & \frac{\partial}{\partial y} \end{bmatrix} \cdot \begin{bmatrix} \frac{\partial}{\partial x} \\ \frac{\partial}{\partial y} \end{bmatrix} \cdot f = \nabla \cdot \nabla f$$

“divergence of gradient”


rotationally invariant
second derivative for 2D functions

<table border="1" style="border-color: red; width: 100px; height: 100px; margin: auto;"> <tr><td>0</td><td>0</td><td>0</td></tr> <tr><td>1</td><td>-2</td><td>1</td></tr> <tr><td>0</td><td>0</td><td>0</td></tr> </table>	0	0	0	1	-2	1	0	0	0	+	<table border="1" style="border-color: red; width: 100px; height: 100px; margin: auto;"> <tr><td>0</td><td>1</td><td>0</td></tr> <tr><td>0</td><td>-2</td><td>0</td></tr> <tr><td>0</td><td>1</td><td>0</td></tr> </table>	0	1	0	0	-2	0	0	1	0	=	<table border="1" style="border-color: red; width: 100px; height: 100px; margin: auto;"> <tr><td>0</td><td>1</td><td>0</td></tr> <tr><td>1</td><td>-4</td><td>1</td></tr> <tr><td>0</td><td>1</td><td>0</td></tr> </table>	0	1	0	1	-4	1	0	1	0
0	0	0																													
1	-2	1																													
0	0	0																													
0	1	0																													
0	-2	0																													
0	1	0																													
0	1	0																													
1	-4	1																													
0	1	0																													
$\underbrace{\hspace{10em}}$		$\underbrace{\hspace{10em}}$																													
<p><i>rate of change for the rate of change in x-direction</i></p> <table style="margin: auto;"> <tr> <td>$\begin{bmatrix} 0 & 0 & 0 \\ 0 & -1 & 1 \\ 0 & 0 & 0 \end{bmatrix}$</td> <td style="font-size: 1.5em;">-</td> <td>$\begin{bmatrix} 0 & 0 & 0 \\ -1 & 1 & 0 \\ 0 & 0 & 0 \end{bmatrix}$</td> </tr> </table>	$\begin{bmatrix} 0 & 0 & 0 \\ 0 & -1 & 1 \\ 0 & 0 & 0 \end{bmatrix}$	-	$\begin{bmatrix} 0 & 0 & 0 \\ -1 & 1 & 0 \\ 0 & 0 & 0 \end{bmatrix}$		<p><i>rate of change for the rate of change in y-direction</i></p> <table style="margin: auto;"> <tr> <td>$\begin{bmatrix} 0 & 1 & 0 \\ 0 & -1 & 0 \\ 0 & 0 & 0 \end{bmatrix}$</td> <td style="font-size: 1.5em;">-</td> <td>$\begin{bmatrix} 0 & 0 & 0 \\ 0 & 1 & 0 \\ 0 & -1 & 0 \end{bmatrix}$</td> </tr> </table>	$\begin{bmatrix} 0 & 1 & 0 \\ 0 & -1 & 0 \\ 0 & 0 & 0 \end{bmatrix}$	-	$\begin{bmatrix} 0 & 0 & 0 \\ 0 & 1 & 0 \\ 0 & -1 & 0 \end{bmatrix}$																							
$\begin{bmatrix} 0 & 0 & 0 \\ 0 & -1 & 1 \\ 0 & 0 & 0 \end{bmatrix}$	-	$\begin{bmatrix} 0 & 0 & 0 \\ -1 & 1 & 0 \\ 0 & 0 & 0 \end{bmatrix}$																													
$\begin{bmatrix} 0 & 1 & 0 \\ 0 & -1 & 0 \\ 0 & 0 & 0 \end{bmatrix}$	-	$\begin{bmatrix} 0 & 0 & 0 \\ 0 & 1 & 0 \\ 0 & -1 & 0 \end{bmatrix}$																													
$\frac{\partial^2 f}{\partial x^2} = \frac{\partial f}{\partial x} \left(\frac{+1}{2} \right) - \frac{\partial f}{\partial x} \left(\frac{-1}{2} \right)$		$\frac{\partial^2 f}{\partial y^2} = \frac{\partial f}{\partial y} \left(\frac{+1}{2} \right) - \frac{\partial f}{\partial y} \left(\frac{-1}{2} \right)$																													

Laplacian of a Gaussian (LoG)

$\Delta * G$ image should
be smoothed a bit first

$$LoG(x, y) = -\frac{1}{\pi\sigma^4} \left[1 - \frac{x^2 + y^2}{2\sigma^2} \right] \cdot e^{-\frac{x^2 + y^2}{2\sigma^2}}$$



MATLAB: `logfilt = fspecial('log',25,4);`


Second Image Derivatives

(Laplace operator Δf)

$$\Delta f = \frac{\partial^2 f}{\partial x^2} + \frac{\partial^2 f}{\partial y^2}$$




For simplicity, assume $f(x,y) = \text{const}(y)$.
 Then, Laplacian of f is simply
 a second derivative of $f(x) = f(x,y)$

Application: Laplacian Zero Crossings are used for edge detection
 (alternative to methods computing Gradient extrema)



Laplacian of a Gaussian (LoG)

$\Delta * G$ image should
be smoothed a bit first



Unsharp masking

- What does blurring take away?

$$U = I - G * I$$



-



=

+ α 


=



Unsharp masking

$$(1 + \alpha)I - \alpha \cdot G * I \approx [(1 + \alpha)G_{\sigma_1} - \alpha \cdot G_{\sigma_2}] * I$$

$$\sigma_1 \ll \sigma_2$$



$$U = I - G * I$$

$$I + \alpha \cdot U$$



+α



=





Unsharp masking

MATLAB

```

Imrgb = imread('file.jpg');
im = im2double(rgb2gray(imrgb));
g= fspecial('gaussian', 25,4);
imblur = conv2(im,g,'same');
imagesc([im imblur])
imagesc([im im+.4*(im-imblur)])
  
```



<http://homepages.inf.ed.ac.uk/rbf/HIPR2/unsharp.htm>



Reading: Forsyth & Ponce ch.7.5

Filters and Templates

- Applying a filter at some point can be seen as taking a dot-product between the image and some vector
- Filtering the image is a set of dot products
- Insight
 - filters may look like the effects they are intended to find
 - filters find effects they look like





Normalized Cross-Correlation (NCC)

- filtering as a **dot product**
- now **measure the angle**:

NCC output is **filter output**
divided by root of the sum of
squares of values over which
filter lies

h

template (filter, kernel, mask)
of size $n = (2k+1) \times (2k+1)$



$h \cdot f_t = \sum_{i=1}^n h[i] f_t[i]$

dot product

cross-correlation of h and f at $t=(x,y)$

$$\frac{\sum_{u=-k}^k \sum_{v=-k}^k h[u,v] \cdot f[x+u, y+v]}{|h| \cdot |f_t|}$$

division makes this
a non-linear operation

$$g[t] = \frac{h \cdot f_t}{|h| \cdot |f_t|} = \cos(\alpha)$$

vector lengths $|z| = \sqrt{\sum_{i=1}^n z_i^2}$



Normalized Cross-Correlation (NCC)

- filtering as a **dot product**
- now **measure the angle:**

NCC output is **filter output**
divided by **root of the sum of**
squares of values over which
filter lies

h

template (filter, kernel, mask)
of size $n = (2k+1) \times (2k+1)$



these vectors do not have to be
in the “positive” quadrant

■ Tricks:

- subtract *template average* \bar{h}
(to give zero output for constant
regions, reduces response to
irrelevant background)
- subtract *patch average* \bar{f}_t when
computing the normalizing
constant (i.e. subtract the image
mean in the neighborhood)

$$g[t] = \frac{(h - \bar{h}) \cdot (f_t - \bar{f}_t)}{|h - \bar{h}| \cdot |f_t - \bar{f}_t|}$$

NCC

Normalized Cross-Correlation (NCC)

- filtering as a **dot product**

- now **measure the angle:**



NCC output is **filter output**
divided by root of the sum of
squares of values over which
filter lies

Tricks:

- subtract *template average* \bar{h}
(to give zero output for constant
regions, reduces response to
irrelevant background)
- subtract *patch average* \bar{f}_t when
computing the normalizing
constant (i.e. subtract the image
mean in the neighborhood)

h

template (filter, kernel, mask)
of size $n = (2k+1) \times (2k+1)$



these vectors do not have to be
in the “positive” quadrant

equivalently using statistical term σ (standard deviation)

$$g[t] = \frac{(h - \bar{h}) \cdot (f_t - \bar{f}_t)}{n \cdot \sigma_h \cdot \sigma_{f_t}} = \text{NCC} = \text{cov}(h, f_t)$$

Remember: st.div. $\sigma_z \equiv \sqrt{\frac{1}{n} \sum_{i=1}^n (z_i - \bar{z})^2} = \sqrt{\frac{1}{n}} \cdot |z - \bar{z}|$

Normalized Cross-Correlation (NCC)


- filtering as a **dot product**
- now **measure the angle:**
NCC output is **filter output**
divided by root of the sum of
squares of values over which
filter lies

■ Tricks:

- subtract *template average* \bar{h}
(to give zero output for constant regions, reduces response to irrelevant background)
- subtract *patch average* \bar{f}_t when computing the normalizing constant (i.e. subtract the image mean in the neighborhood)


 h

template (filter, kernel, mask)
of size $n = (2k+1) \times (2k+1)$



standard in statistics
correlation coefficient

ρ \longleftrightarrow
between h and f_t


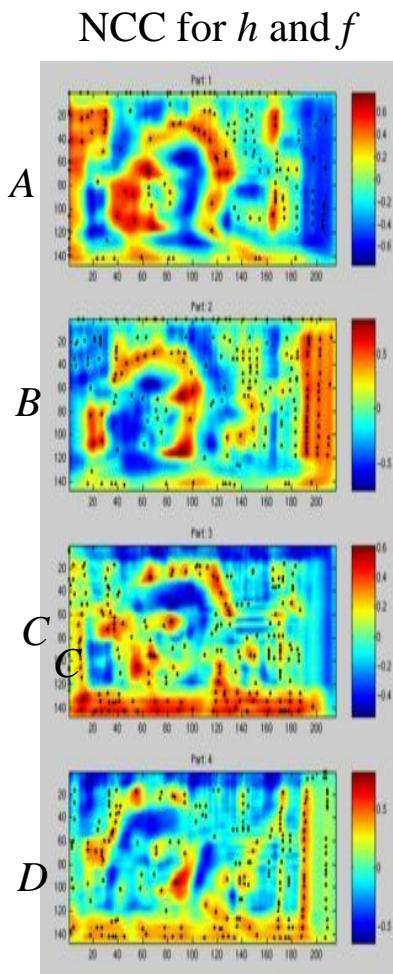
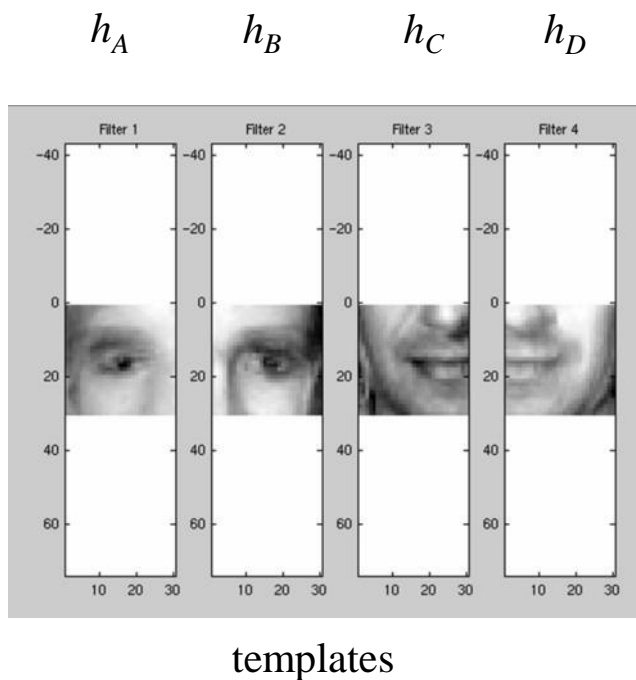
equivalently using statistical term cov (covariance)

$$g[t] = \frac{\text{cov}(h, f_t)}{\sigma_h \cdot \sigma_{f_t}}$$

NCC

$$\text{cov}(a, b) \equiv E(a - \bar{a})(b - \bar{b}) = \frac{1}{n} \sum_{i=1}^n (a_i - \bar{a})(b_i - \bar{b}) = \frac{(a - \bar{a}) \cdot (b - \bar{b})}{n}$$

Normalized Cross-Correlation (NCC)



points mark local maxima of NCC
for each template


points of interest or **feature points**

Other features... (Szeliski sec 4.1.1)

- Feature points are used for:
 - Image alignment (homography, fundamental matrix)
 - 3D reconstruction
 - Motion tracking
 - Object recognition
 - Indexing and database retrieval
 - Robot navigation
 - ... other


Harris corner detector

- C.Harris, M.Stephens. “A Combined Corner and Edge Detector”. 1988




The Basic Idea


- We should easily recognize the point by looking through a small window
- Shifting a window in *any direction* should give *a large change* in intensity




Harris Detector: Basic Idea



“flat” region:
no change in all
directions



“edge”:
no change along the edge
direction




“corner”:
significant change in all
directions

Harris Detector: Mathematics

For any given image patch or window w
we should measure how it changes
when shifted by $ds = \begin{bmatrix} u \\ v \end{bmatrix}$

Notation: let patch be defined
by its support function $w(x,y)$
over image pixels



Harris Detector: Mathematics

patch w change measure for shift $ds = \begin{bmatrix} u \\ v \end{bmatrix}$: weighted sum of squared differences


$$E_w(u, v) := \sum_{x, y} w(x, y) \cdot [I(x+u, y+v) - I(x, y)]^2$$

Window
function


Shifted
intensity

Intensity

NOTE:
window support
functions $w(x, y) =$



1 in window, 0 outside



Gaussian
(weighted) support

Harris Detector: Mathematics

Change of intensity for the **shift** $ds = \begin{bmatrix} u \\ v \end{bmatrix}$ assuming **image gradient** $\nabla I \equiv \begin{bmatrix} I_x \\ I_y \end{bmatrix}$

$$I(x+u, y+v) - I(x, y) \approx I_x \cdot u + I_y \cdot v = ds^T \cdot \nabla I$$

rate of change for I at (x, y) in direction $(u, v) = ds$

(remember **gradient** definition on earlier slides!!!!)

this is 2D analogue of 1st order Taylor expansion

$$[I(x+u, y+v) - I(x, y)]^2 \approx ds^T \cdot \nabla I \cdot \nabla I^T \cdot ds$$

$$E_w(u, v) = \sum_{x, y} w(x, y) \cdot [I(x+u, y+v) - I(x, y)]^2$$

$$\approx ds^T \cdot \left(\sum_{x, y} w(x, y) \cdot \nabla I \cdot \nabla I^T \right) \cdot ds = ds^T \cdot M_w \cdot ds$$

↖ M_w

Harris Detector: Mathematics

Change of intensity for the **shift** $ds = \begin{bmatrix} u \\ v \end{bmatrix}$ assuming **image gradient** $\nabla I \equiv \begin{bmatrix} I_x \\ I_y \end{bmatrix}$

$$E_w(u, v) \cong [u \ v] \cdot M_w \cdot \begin{bmatrix} u \\ v \end{bmatrix} = ds^T \cdot M_w \cdot ds$$

where M_w is a 2×2 matrix computed from image derivatives inside patch w

matrix M is also called
Harris matrix or *structure tensor*

$$\dots \cdot \left(\sum_{x, y} w(x, y) \cdot \overbrace{\nabla I \cdot \nabla I^T}^{M_w} \right) \cdot \dots$$


This tells you how
to compute M_w
at any window w
(t.e. any image patch)

Harris Detector: Mathematics

Change of intensity for the **shift** $ds = \begin{bmatrix} u \\ v \end{bmatrix}$ assuming **image gradient** $\nabla I \equiv \begin{bmatrix} I_x \\ I_y \end{bmatrix}$


$$E_w(u, v) \cong [u \ v] \cdot M_w \cdot \begin{bmatrix} u \\ v \end{bmatrix} = ds^T \cdot M_w \cdot ds$$

paraboloid



M is a positive semi-definite matrix (**Exercise:** show that $ds^T \cdot M \cdot ds \geq 0$ for any ds)

M can be analyzed via its *isolines*, e.g. $ds^T \cdot M_w \cdot ds = 1$ (ellipsoid)




two eigen values of matrix M_w

Points on this ellipsoid are shifts $ds=(u,v)$ giving the same value of energy $E(u,v)=1$. Thus, the ellipsoid allows to visually compare sensitivity of energy E to shifts ds in different directions

Harris Detector: Mathematics

Classification of image points
using eigenvalues of M :



Harris Detector: Mathematics

Measure of corner response:

$$R = \frac{\det M}{\text{Trace } M}$$

$$\det M = \lambda_1 \lambda_2$$

$$\text{trace } M = \lambda_1 + \lambda_2$$

R should be large

(it implies that both λ are far from zero)

Harris Detector

■ The Algorithm:


- Find points with large corner response function R
 $R > \text{threshold}$
- Take the points of local maxima of R

Harris Detector: Workflow




Harris Detector: Workflow

Compute corner response R




Harris Detector: Workflow

Find points with large corner response: $R > \text{threshold}$



Harris Detector: Workflow

Take only the points of local maxima of R




Harris Detector: Workflow



Harris Detector: Some Properties

■ Rotation invariance



Ellipse rotates but its shape (i.e. eigenvalues) remains the same



Corner response R is invariant to image rotation

Harris Detector: Some Properties

■ Partial invariance to *affine intensity* change

✓ Only derivatives are used \Rightarrow invariance to intensity shift $I \rightarrow I + b$


✓ Intensity scale: $I \rightarrow a I$



features locations stay the same,
but some may appear or disappear depending on gain a

Harris Detector: Some Properties

- But: non-invariant to *image scale*!




All points will be
classified as **edges**

Corner !


Scale Invariant Detection

- Consider regions (e.g. circles) of different sizes around a point
- Regions of corresponding sizes will look the same in both images



Scale Invariant Detection

- The problem: how do we choose corresponding circles *independently* in each image?
- Choose the scale of the “best” corner



Other feature detectors

- *LoG* and *DoG* operators are also used to detect “features”
- they find reliable “blob” features (at appropriate scale)
- these operators also respond to edges. To improve “selectivity”, post-processing is necessary.
 - e.g. eigen-values of the Harris matrix could be used as in the corner operator.
If the ratio of the eigen-values is too high, then the local image is regarded as too edge-like and the feature is rejected.

Other features

■ MOPS, Hog, SIFT, ...

Features are characterized by **location** and **descriptor**

color

any pixel

RGB vector

edge

Laplacian zero crossing

image gradient

corner

local max of R magnitude of R

MOPS

corners

normalized intensity patch


more below

HOG
SIFTLOG extrema points
or other interest pointsgradient orientation
histogramshighly
discriminative

(see Szeliski, Sec. 4.1.2)

Feature descriptors

- We know how to detect points
- Next question: **How to match them?**




Point descriptor should be:

1. Invariant
2. Distinctive

Descriptors Invariant to Rotation

- Find local orientation

Dominant direction of gradient



- Extract image patches relative to this orientation

Multi-Scale Oriented Patches (MOPS)

- Interest points
 - Multi-scale Harris corners
 - Orientation from blurred gradient
 - Geometrically invariant to rotation
- Descriptor vector
 - Bias/gain normalized sampling of local patch (8x8)
 - Photometrically invariant to affine changes in intensity

[Brown, Szeliski, Winder, CVPR'2005]

Descriptor Vector

- Orientation = blurred gradient
- Rotation Invariant Frame
 - Scale-space position (x, y, s) + orientation (θ)



Detections at multiple scales




Figure 1. Multi-scale Oriented Patches (MOPS) extracted at five pyramid levels from one of the Matier images. The boxes show the feature orientation and the region from which the descriptor vector is sampled.

MOPS descriptor vector

- 8x8 oriented patch
 - Sampled at 5 x scale
- Bias/gain normalisation: $I' = (I - \mu) / \sigma$

